

Cooperation and Competition in Multi-Agent Reinforcement Learning (via Pac-Man)

Motivation

- Multi-Agent systems can be used to effectively address problems in domains such as robotics, distributed control, telecommunications and economics.
- The complexity of these tasks makes it difficult to solve with preprogrammed agent behaviours. The agents must instead discover a solution via learning.
- Many tasks can be dealt with in a cooperative manner. It is important that we are able to train agents to act cooperatively by sharing instantaneous information, episodic experience and learned knowledge.
- We attempt to simulate state-of-the-art algorithms in MARL with close attention to cooperative learning algorithms and Self Play algorithms for training.

Problem Formulation

- We will be using the Pac-Man API from CS 188 of University of California, Berkely.
- We put two opposing teams in a maze-like grid (as typically seen in Pac-Man games).
- For the simplest version of this problem:
 - The team consists of two Pac-Mans.
 - The grid is symmetrically divided into areas for each team.
 - The objective of each team is to maximize its score before terminal conditions are met and the game ends.
 - This can be done by eating the coins present in opponent areas.
 - This can be done by eating opponent Pac-Mans in your territory.
 - There is also a negative living reward to motivate the agents.
 - Terminal conditions:
 - The clock time runs out
 - There is only one Pac-Man team that has not been eaten
 - Training is done in a Self Play manner, where a team plays against itself to iteratively improve its policy.
- We can arbitrarily increase the complexity of the game to enforce greater cooperation between the agents and see interesting behaviors:
 - Information sharing can only be done when Pac-Mans are in adjacent spaces
 - We limit their vision upto a certain number of spaces.
 - A team can share what each member is currently seeing.
 - We increase the number of teams and agents per team.
 - We can add a benign extra team, that agents can learn to cooperate with or take advantage of.




Fig. Screenshot from Contest-2 of CS 188 Fall 2018 [Link](#)

Proposed Learning Algorithm

- State-Of-The-Art
 - PPO ([Link](#))
 - Multi-Agent MDPs - Multiple Actors, Centralized Critics
 - LOLA ([Link](#))

Existing Work

- **Iterative improvement using adversarial agents via Self-Play ([Link](#)):** AlphaGo Zero has learned by playing against itself. According to its creators, it surpassed the performance of AlphaGo which was trained by playing against professionals.
- **Emergent Behaviors in Multi-Agent Interaction ([Link](#)):** Here they have explored the Cooperative and Adversarial training of agents in a setup of Hiding and Seek game where the two teams with different objectives(one team has to hide and the other has to seek) discover newer and unexpected strategies.
- **Cooperative Agents ([Link](#)):** In the GitHub link for Books and Papers we have a Review paper with multiple works on Cooperative Multi-Agent Learning.

Baseline Algorithms

- SARSA, Q-Learning (Classical/Deep Learning)
- MC, TD(n)
- Policy Gradient
- Actor-Critic Methods

References:

- Cooperative Agents: ([Link 1](#) , [Link 2](#) , [Link 3](#))
- Competition, Sharing and Self Play: ([Link](#))